

HW 4 Solutions

Shravan Vasishth

Exercise 1

```
data("df_pupil_complete")
df_pupil_complete <- df_pupil_complete %>%
  mutate(c_load = load - mean(load))
```

Fit a “maximal” model (correlated varying intercept and slopes for subjects) assuming a normal likelihood.

```
fit_pupil <- brm(p_size ~ 1 + c_load + (c_load | subj),
  data = df_pupil_complete,
  family = gaussian(),
  prior = c(
    prior(normal(1000, 500), class = Intercept),
    prior(normal(0, 1000), class = sigma),
    prior(normal(0, 100), class = b, coef = c_load),
    prior(normal(0, 1000), class = sd),
    prior(lkj(2), class = cor)),
  control=list(adapt_delta=0.99, max_treedepth=15))
```

(a) Examine the effect of load on pupil size, and the average pupil size.

```
fit_pupil

## Family: gaussian
## Links: mu = identity; sigma = identity
## Formula: p_size ~ 1 + c_load + (c_load | subj)
## Data: df_pupil_complete (Number of observations: 2228)
## Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
##        total post-warmup draws = 4000
##
## Multilevel Hyperparameters:
## ~subj (Number of levels: 20)
##                               Estimate Est.Error l-95% CI u-95% CI Rhat
## sd(Intercept)            3355.39    449.51  2546.78  4316.00 1.00
## sd(c_load)                71.45     15.23    47.53   106.55 1.00
## cor(Intercept,c_load)    0.30      0.24    -0.23     0.71 1.01
##                               Bulk_ESS Tail_ESS
## sd(Intercept)              717       1008
## sd(c_load)                 1267      1951
## cor(Intercept,c_load)     1415      1886
##
## Regression Coefficients:
##                               Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept  2469.01    481.63  1533.89  3391.77 1.00       638      952
```

```

## c_load      40.22     24.87    -9.67     88.54 1.01      1112      1580
##
## Further Distributional Parameters:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sigma    505.09      7.65   490.61   520.07 1.00      4557     2437
##
## Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).

```

If we want to just report the change in pupil size as a function of one unit increase in (centered) load:

```
posterior_summary(fit_pupil, variable = "b_c_load")
```

```

##           Estimate Est.Error Q2.5 Q97.5
## b_c_load     40.2       24.9 -9.67  88.5

```

There isn't a clear effect of load across all the subjects.

But the intercept seems to be quite large,

```
posterior_summary(fit_pupil, variable = "b_Intercept")
```

```

##           Estimate Est.Error Q2.5 Q97.5
## b_Intercept 2469        482 1534  3392

```

and we assumed that it shouldn't be:

```
brms::prior_summary(fit_pupil)
```

```

##           prior class     coef group resp dpar npar lb ub
## (flat)      b
## normal(0, 100) b     c_load
## normal(1000, 500) Intercept
## lkj_corr_cholesky(2) L
## lkj_corr_cholesky(2) L     subj
## normal(0, 1000) sd
## normal(0, 1000) sd     subj
## normal(0, 1000) sd     c_load subj
## normal(0, 1000) sd Intercept subj
## normal(0, 1000) sigma
## source
## default
## user
## user
## user
## (vectorized)
## user
## (vectorized)
## (vectorized)
## (vectorized)
## user

```

See the row with class Intercept.

So maybe our prior for the intercept was overly informative.

(b) Do a sensitivity analysis for the prior on the intercept (α). What is the estimate of the effect (β) under different priors?

We'll try a wider prior for α . A more complete sensitivity analysis would investigate several possible priors

$$\alpha \sim Normal(4000, 2000)$$

```
fit_pupil_2 <- brm(p_size ~ 1 + c_load + (c_load | subj),
  data = df_pupil_complete,
  family = gaussian(),
  prior = c(
    prior(normal(4000, 2000), class = Intercept),
    prior(normal(0, 1000), class = sigma),
    prior(normal(0, 100), class = b, coef = c_load),
    prior(normal(0, 2000), class = sd),
    prior(lkj(2), class = cor))
)

fit_pupil_2

## Family: gaussian
## Links: mu = identity; sigma = identity
## Formula: p_size ~ 1 + c_load + (c_load | subj)
## Data: df_pupil_complete (Number of observations: 2228)
## Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
##         total post-warmup draws = 4000
##
## Multilevel Hyperparameters:
## ~subj (Number of levels: 20)
##           Estimate Est.Error l-95% CI u-95% CI Rhat
## sd(Intercept)     2548.48    399.12   1896.16   3459.68 1.01
## sd(c_load)        69.05     14.58    46.09    103.36 1.00
## cor(Intercept,c_load)  0.27     0.20    -0.17     0.62 1.00
##           Bulk_ESS Tail_ESS
## sd(Intercept)      872      1498
## sd(c_load)        1232      1743
## cor(Intercept,c_load) 1950      2526
##
## Regression Coefficients:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept  5464.64    572.30   4338.88   6579.55 1.01      549      773
## c_load     58.42     17.33    22.88    90.24 1.00     1411     1730
##
## Further Distributional Parameters:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sigma     505.40     7.79   490.45   520.88 1.00      5315     2559
##
## Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).
```

Now there seems to be a clear effect! Our bad prior for the intercept was messing up our inferences!

(c) Is the effect of load consistent across subjects?

```
## For the hierarchical model, this is more complicated,
# because we want the effect (beta) + adjustment:
# we extract the overall group level effect:
beta <- c(as_draws_df(fit_pupil_2)$b_c_load)
# We extract the individual adjustments
ind_effects_v <- paste0("r_subj[", unique(df_pupil_complete$subj), ",c_load]")
adjustment <- as.matrix(as_draws_df(fit_pupil)[ind_effects_v])

## Warning: Dropping 'draws_df' class as required metadata was removed.

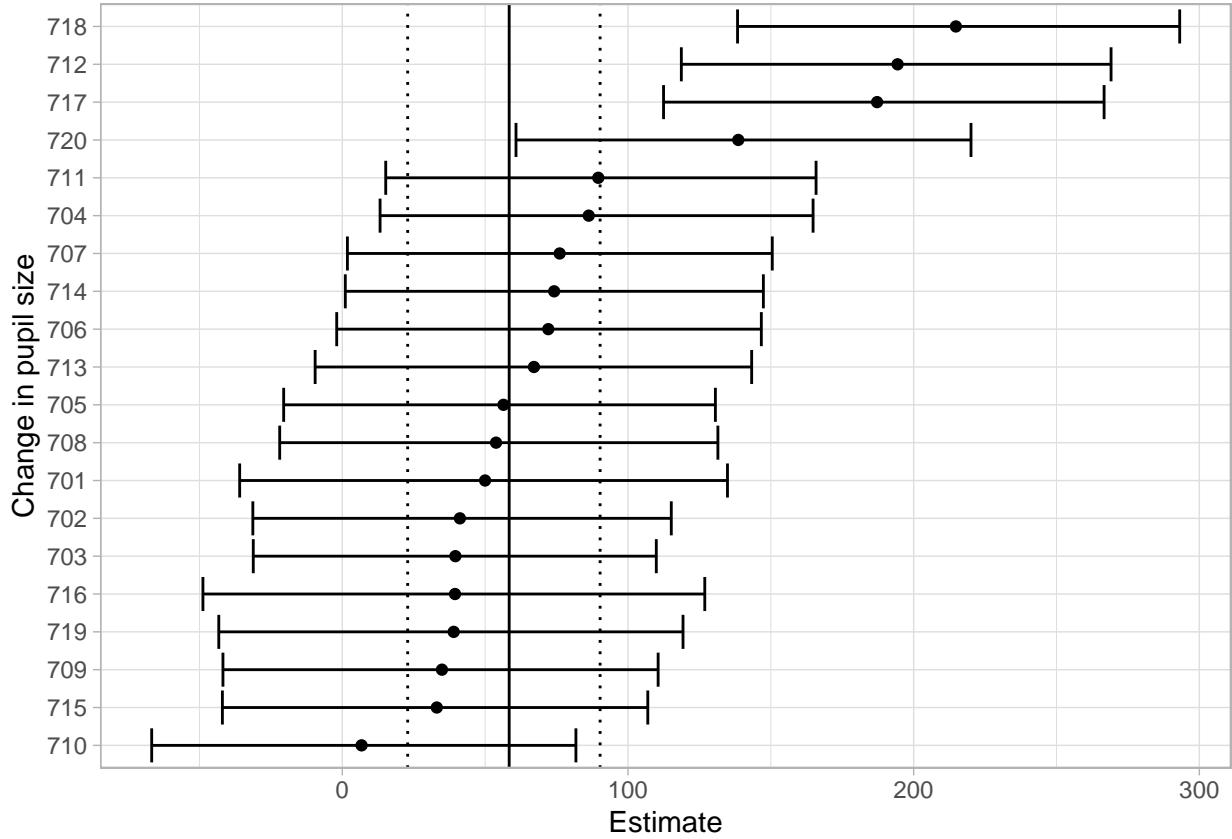
# We get the by subject effects in a data frame where each adjustment
# is in each column.
by_subj_effect <- as_tibble(beta + adjustment)
# We summarize them by getting a table with the mean and the
# quantiles for each column and then binding them.
par_h <- lapply(by_subj_effect, function(x) {
  tibble(
    Estimate = mean(x),
    Q2.5 = quantile(x, .025),
    Q97.5 = quantile(x, .975)
  )
}) %>%
  bind_rows() %>%
  # We add a column to identify that the model,
  # and one with the subject labels:
  mutate(
    subj = unique(df_pupil_complete$subj))%>%
  arrange(Estimate) %>%
  mutate(subj = factor(subj, levels = subj))

ggplot(
  par_h,
  aes(
    ymin = Q2.5, ymax = Q97.5, x = subj, y = Estimate
  )
) +
  geom_errorbar() +
  geom_point() +
  # We'll also add the mean and 95% CrI of the overall difference
  # to the plot:
  geom_hline(
    yintercept =
      posterior_summary(fit_pupil_2)["b_c_load", "Estimate"]
  ) +
  geom_hline(
    yintercept =
      posterior_summary(fit_pupil_2)["b_c_load", "Q2.5"],
    linetype = "dotted", linewidth = 0.5
  ) +
  geom_hline(
    yintercept =
      posterior_summary(fit_pupil_2)["b_c_load", "Q97.5"],
```

```

    linetype = "dotted", linewidth = 0.5
) +
xlab("Change in pupil size") +
coord_flip()

```



Exercise 2

We are going to fit the data with the following log-normal likelihood:

$$rt_n \sim LogNormal(\alpha + u_{i[n],0} + w_{j[n],0} + c_cond_n \cdot (\beta + u_{i[n],1} + w_{j[n],1}), \sigma) \quad (1)$$

We can use similar priors as we used for the Stroop task paying attention to use a not too tight prior for β :

- Priors:

$$\begin{aligned} \alpha &\sim Normal(6, 1.5) \\ \beta &\sim Normal(0, 0.1) \\ \sigma &\sim Normal_+(0, 1) \end{aligned} \quad (2)$$

Here, we will need priors for the group-level parameters. Given that we assume a correlation between by-subject intercept and slope, and by-item intercept and slope, our model has the following structure which requires us to assign priors to Σ_u and Σ_w

$$\begin{aligned} \begin{pmatrix} u_{i,0} \\ u_{i,1} \end{pmatrix} &\sim \mathcal{N} \left(\begin{pmatrix} 0 \\ 0 \end{pmatrix}, \Sigma_u \right) \\ \begin{pmatrix} w_{i,0} \\ w_{i,1} \end{pmatrix} &\sim \mathcal{N} \left(\begin{pmatrix} 0 \\ 0 \end{pmatrix}, \Sigma_w \right) \end{aligned} \quad (3)$$

$$\begin{aligned}\boldsymbol{\Sigma}_u &= \begin{pmatrix} \tau_{u_0}^2 & \rho_u \tau_{u_0} \tau_{u_1} \\ \rho_u \tau_{u_0} \tau_{u_1} & \tau_{u_1}^2 \end{pmatrix} \\ \boldsymbol{\Sigma}_w &= \begin{pmatrix} \tau_{w_0}^2 & \rho_w \tau_{w_0} \tau_{w_1} \\ \rho_w \tau_{w_0} \tau_{w_1} & \tau_{w_1}^2 \end{pmatrix}\end{aligned}\tag{4}$$

In practice this means that we need priors for the by-subject and by-item variances and correlations:

$$\begin{aligned}\tau_{u_0} &\sim Normal_+(0, 1) \\ \tau_{u_1} &\sim Normal_+(0, 1) \\ \rho_u &\sim LKJcorr(2) \\ \tau_{w_0} &\sim Normal_+(0, 1) \\ \tau_{w_1} &\sim Normal_+(0, 1) \\ \rho_w &\sim LKJcorr(2)\end{aligned}\tag{5}$$

Read the data:

```
data("df_gg05_rc")
df_gg05_rc <- df_gg05_rc %>
  mutate(c_cond = if_else(condition == "objgap", 1/2, -1/2))
```

We're ready to fit a model now:

```
fit_df_gg05_rc <- brm(RT ~ c_cond + (c_cond | subj) + (c_cond | item),
  family = lognormal(),
  iter = 10000,
  prior =
  c(
    prior(normal(6, 1.5), class = Intercept),
    prior(normal(0, .1), class = b),
    prior(normal(0, 1), class = sigma),
    prior(normal(0, 1), class = sd),
    prior(lkj(2), class = cor)
  ),
  data = df_gg05_rc,
  control=list(adapt_delta=0.99, max_treedepth=15)
)
```

(a) Examine the effect of relative clause attachment site (the predictor `c_cond`) on reading times RT in log-scale.

We'll focus on β :

```
posterior_summary(fit_df_gg05_rc, variable = "b_c_cond")
```

```
##             Estimate Est.Error   Q2.5 Q97.5
## b_c_cond     0.0984     0.046 0.00444 0.187
```

(b) Estimate the median difference between relative clause attachment sites in milliseconds, and report the mean and 95% CI.

```
alpha <- as_draws_df(fit_df_gg05_rc)$b_Intercept
beta <- as_draws_df(fit_df_gg05_rc)$b_c_cond
# Difference between object RC coded as .5 and subject RC coded as .5
```

```

effect <- exp(alpha + beta * .5) - exp(alpha + beta * -.5)
c(mean = mean(effect), quantile(effect, c(.025,.975)))

##  mean 2.5% 97.5%
## 35.33 1.54 69.35

```

(c) Do a sensitivity analysis. What is the estimate of the effect (β) under different priors? What is the difference in milliseconds between conditions under different priors?

Next, we do a sensitivity analysis using a tighter prior for β , $\beta \sim Normal(0, 0.01)$:

```

fit_df_gg05_rc2 <- brm(RT ~ c_cond + (c_cond | subj) + (c_cond | item),
  family = lognormal(),
  prior =
  c(
    prior(normal(6, 1.5), class = Intercept),
    prior(normal(0, 0.01), class = b),
    prior(normal(0, 1), class = sigma),
    prior(normal(0, 1), class = sd),
    prior(lkj(2), class = cor)
  ),
  data = df_gg05_rc,
  control=list(adapt_delta=0.99, max_treedepth=15)
)

posterior_summary(fit_df_gg05_rc2, variable = "b_c_cond")

```

```

##           Estimate Est.Error   Q2.5   Q97.5
## b_c_cond  0.00425  0.00973 -0.0147  0.0233

```

And here we use the prior, $\beta \sim Normal(0, 1)$:

```

fit_df_gg05_rc3 <- brm(RT ~ c_cond + (c_cond | subj) + (c_cond | item),
  family = lognormal(),
  prior =
  c(
    prior(normal(6, 1.5), class = Intercept),
    prior(normal(0, 1), class = b),
    prior(normal(0, 1), class = sigma),
    prior(normal(0, 1), class = sd),
    prior(lkj(2), class = cor)
  ),
  data = df_gg05_rc,
  control=list(adapt_delta=0.99, max_treedepth=15)
)

posterior_summary(fit_df_gg05_rc3, variable = "b_c_cond")

```

```

##           Estimate Est.Error   Q2.5   Q97.5
## b_c_cond     0.122    0.0511  0.0197  0.22

```

We can summarize the estimates of β given different priors in the following way:

Estimate	Q2.5	Q97.5	Prior for β
0.098	0.004	0.187	Normal(0,01)

Estimate	Q2.5	Q97.5	Prior for β
0.004	-0.015	0.023	Normal(0,.1)
0.122	0.020	0.220	Normal(0,1)

We can take a look at the estimate for the difference between conditions in milliseconds:

```
alpha <- as_draws_df(fit_df_gg05_rc)$b_Intercept
beta1 <- as_draws_df(fit_df_gg05_rc)$b_c_cond
diff1 <- exp(alpha + beta1/2) - exp(alpha - beta1/2)
c(mean = mean(diff1), quantile(diff1, c(.025, .975)))

##  mean 2.5% 97.5%
## 35.33 1.54 69.35
```

If we repeat this with each model, we can take a look at the effect of the prior on the difference between conditions:

Estimate (ms)	Q2.5	Q97.5	Prior for β
35.33	1.53	69.35	Normal(0,.01)
1.47	-5.07	8.12	Normal(0,.1)
44.05	6.72	82.59	Normal(0,1)

This sensitivity shows us that the posterior changes quite a lot under different priors; depending on what we consider to be a reasonable prior (i.e., depending on our prior beliefs), the inference that we derive from the model and data can vary quite a bit when the data are relatively sparse, as in this case.

Exercise 3

```
data("df_gibsonwu")
df_gibsonwu$cond<-ifelse(df_gibsonwu$type=="obj-ext",0.5,-0.5)

data("df_gibsonwu2")
df_gibsonwu2$cond<-ifelse(df_gibsonwu2$condition=="obj-ext",0.5,-0.5)
```

Analysis of Gibson and Wu 2013 data:

Normal likelihood:

```
priorsgwfullnormal <- c(set_prior("normal(500, 200)", class = "Intercept"),
set_prior("normal(0,500)", class = "b",
coef = "cond"),
set_prior("normal(0, 500)", class = "sd"),
set_prior("normal(0, 1000)", class = "sigma"),
set_prior("lkj(2)", class = "cor"))

m_gwfullnormal<-brm(rt~cond + (1+cond|subj)+(1+cond|item),
family=gaussian(), prior=priorsgwfullnormal,
warmup=1000,
iter=2000,
cores=4,
data = df_gibsonwu)
```

Log-normal likelihood:

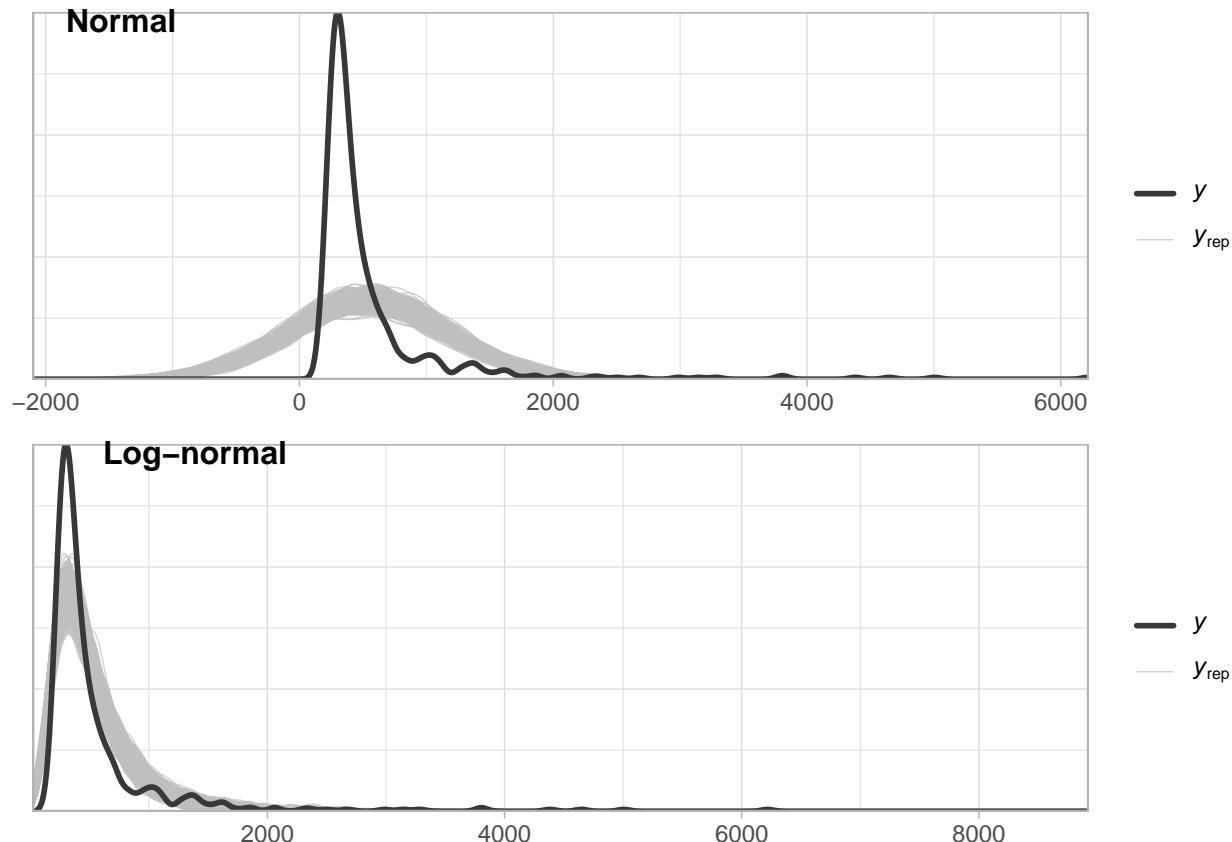
```
priorsgwfulllogn <- c(set_prior("normal(6, 0.6)", class = "Intercept"),
  set_prior("normal(0,1)", class = "b",
            coef = "cond"),
  set_prior("normal(0, 1)", class = "sd"),
  set_prior("normal(0, 1)", class = "sigma"),
  set_prior("lkj(2)", class = "cor"))

m_gwfulllogn<-brm(rt~cond + (1+cond|subj)+(1+cond|item),
  family=lognormal(), prior=priorsgwfulllogn,
  warmup=1000,
  iter=2000,
  cores=4,
  data = df_gibsonwu,
  control=list(adapt_delta=0.99, max_treedepth=15))
```

(a) Posterior predictive distributions

Gibson and Wu (2013) data:

```
pp_gwnormal<-pp_check(m_gwfullnormal,ndraws=1000)
pp_gwlogn<-pp_check(m_gwfulllogn,ndraws=1000)
cowplot::plot_grid(pp_gwnormal,pp_gwlogn,labels=c("Normal","Log-normal"),label_size=12,ncol=1)
```



In the Gibson and Wu (2013) data, the normal likelihood leads to a pretty dramatic mismatch between the observed and posterior predictive values; negative values are generated. The log-normal has a better match.

The replication attempt of Gibson and Wu (2013):

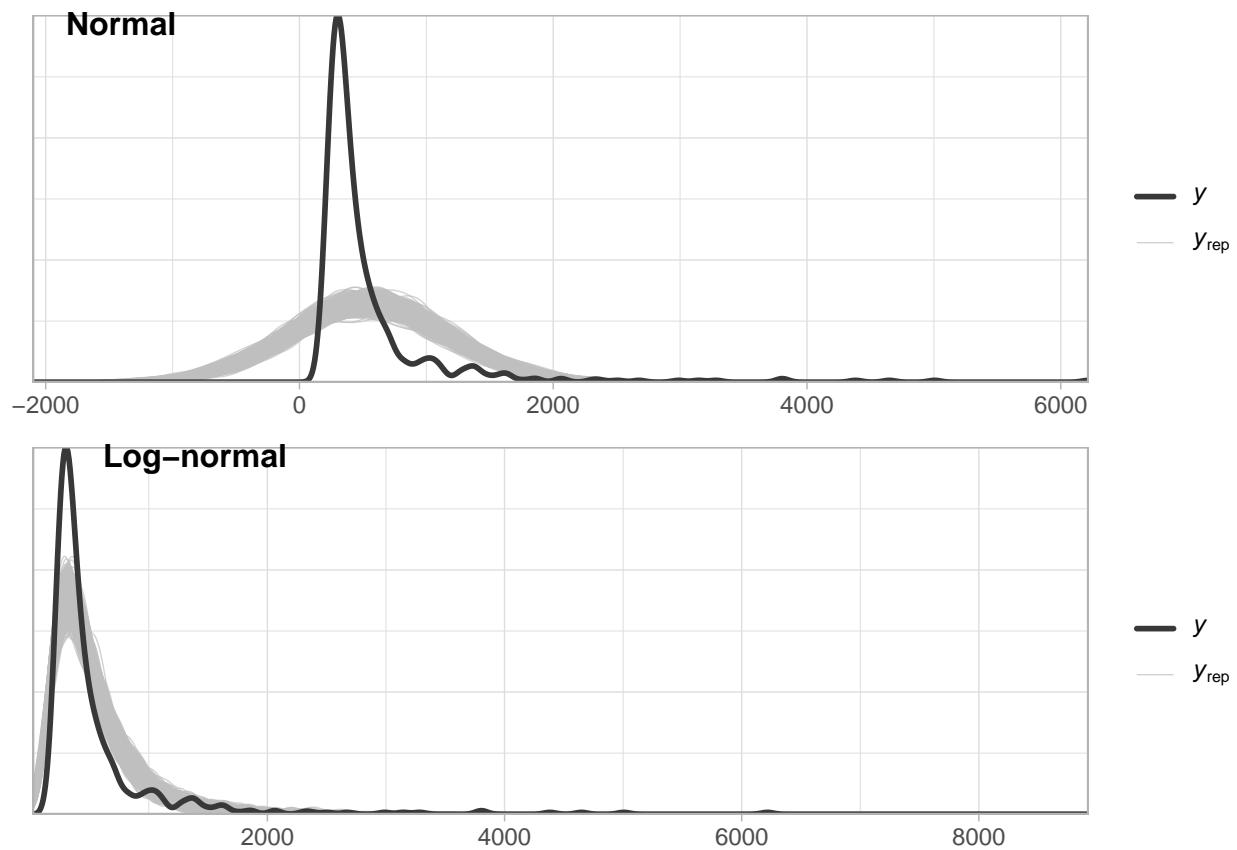
```

m_gw2fullnormal<-brm(rt~cond + (1+cond|subj)+(1+cond|item),
  family=gaussian(), prior=priorsgwfullnormal,
  warmup=1000,
  iter=2000,
  cores=4,
  data = df_gibsonwu2)

m_gw2fulllogn<-brm(rt~cond + (1+cond|subj)+(1+cond|item),
  family=lognormal(), prior=priorsgwfulllogn,
  warmup=1000,
  iter=2000,
  cores=4,
  data = df_gibsonwu2)

pp_gw2normal<-pp_check(m_gw2fullnormal,ndraws=1000)
pp_gw2logn<-pp_check(m_gw2fulllogn,ndraws=1000)
cowplot::plot_grid(pp_gwnormal,pp_gwlogn,labels=c("Normal","Log-normal"),label_size=12,ncol=1)

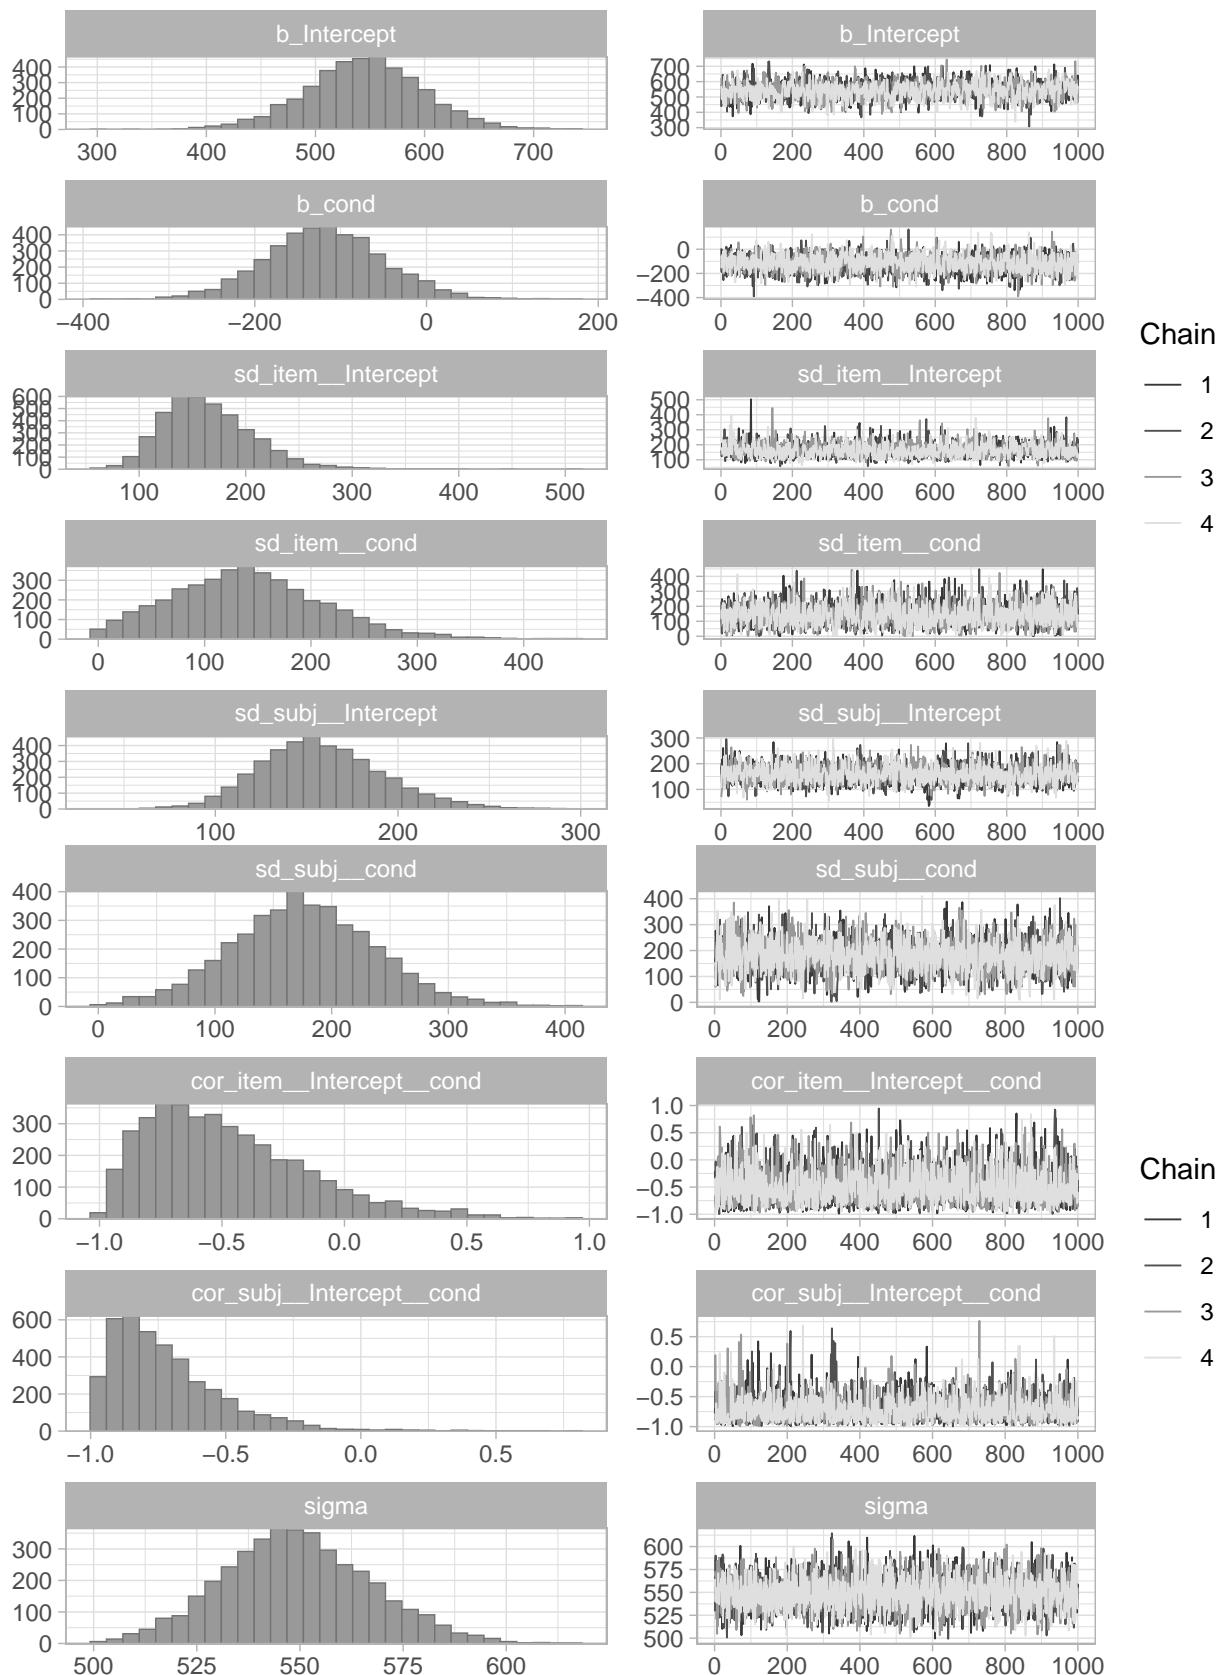
```



Here we have the same situation as in the first data set: assuming a normal likelihood yields a posterior predictive distribution that has negative values; the log-normal produces more realistic values.

(b) Gibson and Wu (2013) effect estimates:

```
plot(m_gwfullnormal)
```



```

gw_intercept <- as_draws_df(m_gwfullnormal)$b_Intercept
gw_slope <- as_draws_df(m_gwfullnormal)$b_cond

gw_interceptln <- as_draws_df(m_gwfulllogn)$b_Intercept
gw_slopeln <- as_draws_df(m_gwfulllogn)$b_cond

gw_RT_diffnormal <- gw_slope
quantile(gw_RT_diffnormal,prob=c(0.025,0.975))

##    2.5% 97.5%
## -255.2   18.6

gw_RT_difflogn <- exp(gw_interceptln + gw_slopeln/2) -
  exp(gw_interceptln - gw_slopeln/2)

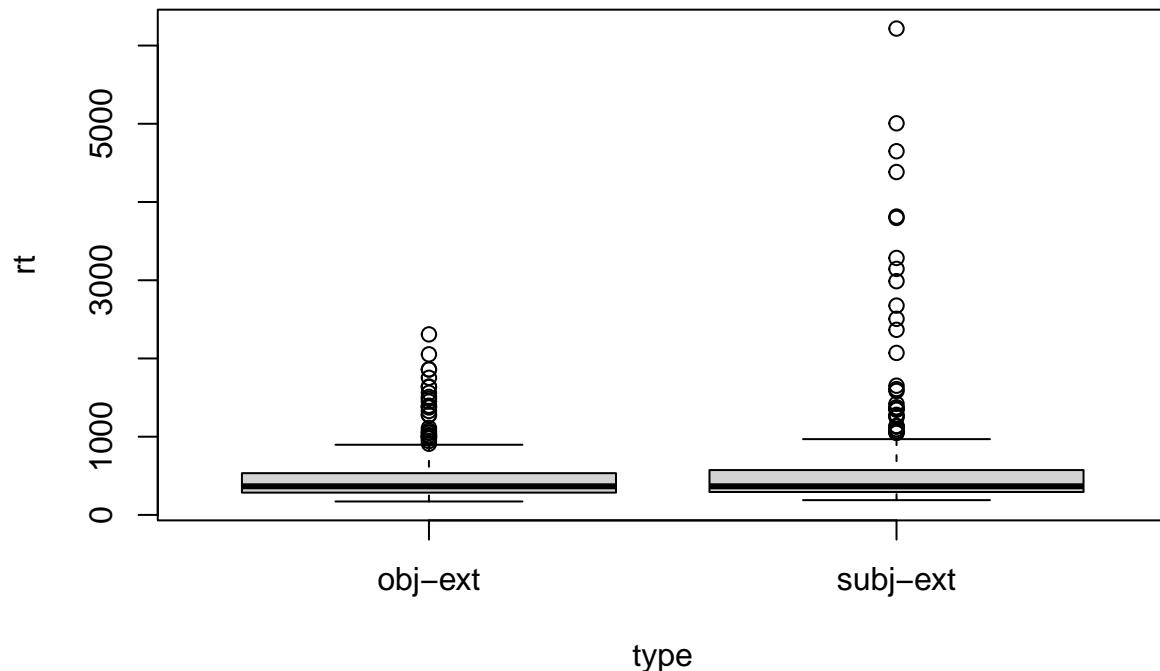
quantile(gw_RT_difflogn,prob=c(0.025,0.975))

##    2.5% 97.5%
## -81.0   16.4

```

The effect estimate is much smaller, and has a smaller 95% credible interval, with the log-normal likelihood. This is because there are a few extreme values in the data that bias the mean effect in subject relatives:

```
boxplot(rt~type,df_gibsonwu)
```



The same thing happens with the replication attempt of the Gibson and Wu design:

```

gw2_intercept <- as_draws_df(m_gw2fullnormal)$b_Intercept
gw2_slope <- as_draws_df(m_gw2fullnormal)$b_cond

gw2_interceptln <- as_draws_df(m_gw2fulllogn)$b_Intercept
gw2_slopeln <- as_draws_df(m_gw2fulllogn)$b_cond

gw2_RT_diffnormal <- gw2_slope

```

```

quantile(gw2_RT_diffnormal,prob=c(0.025,0.975))

##    2.5% 97.5%
## -264.1   53.9

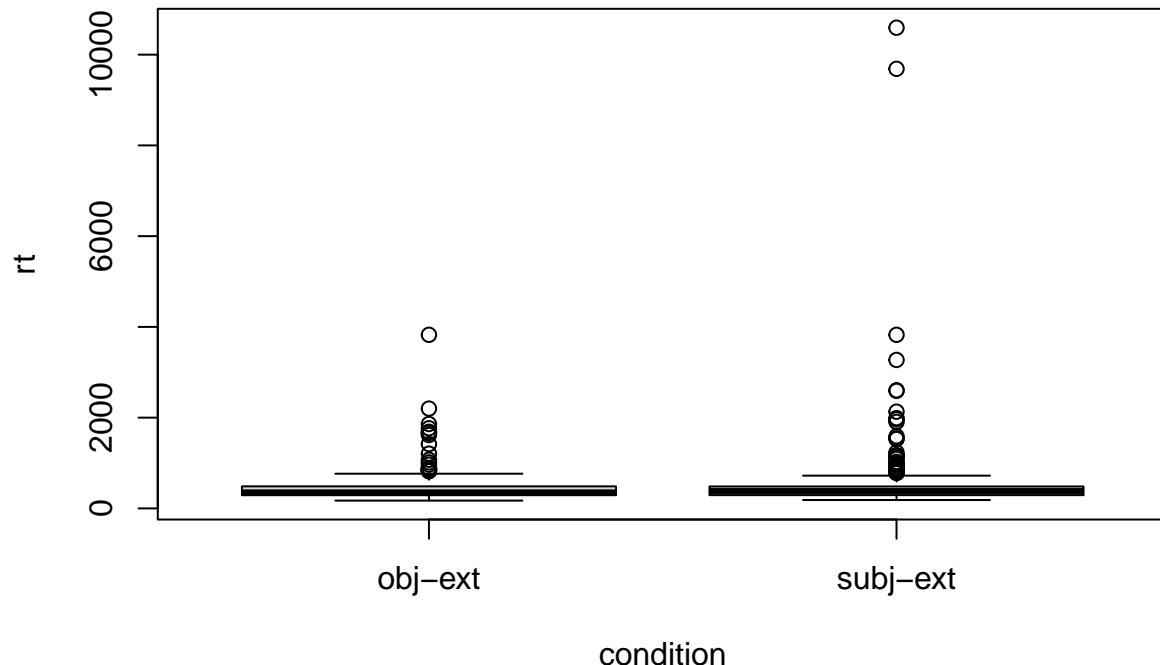
gw2_RT_difflogn <- exp(gw2_interceptln + gw2_slopln/2) -
  exp(gw2_interceptln - gw2_slopln/2)

quantile(gw2_RT_difflogn,prob=c(0.025,0.975))

##    2.5% 97.5%
## -90.5   21.0

boxplot(rt~condition,df_gibsonwu2)

```



- (c) The log-normal is better because it yields more realistic posterior predictive data and the estimate of the effect is not unduly affected by a few extreme values.

Using an informative prior (the posterior of a preceding study) for the next study:

```

summary(m_gw2fulllogn)

## Family: lognormal
## Links: mu = identity; sigma = identity
## Formula: rt ~ cond + (1 + cond | subj) + (1 + cond | item)
## Data: df_gibsonwu2 (Number of observations: 595)
## Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
##        total post-warmup draws = 4000
##
## Multilevel Hyperparameters:
## ~item (Number of levels: 15)
##                               Estimate Est.Error 1-95% CI u-95% CI Rhat
## sd(Intercept)            0.16     0.04    0.10    0.26 1.00
## sd(cond)                 0.19     0.07    0.08    0.34 1.00
## cor(Intercept,cond)     -0.41     0.28   -0.86    0.22 1.00

```

```

##                                Bulk_ESS Tail_ESS
## sd(Intercept)            1347     1493
## sd(cond)                 1157     1712
## cor(Intercept,cond)      2131     2252
##
## ~subj (Number of levels: 40)
##           Estimate Est.Error l-95% CI u-95% CI Rhat
## sd(Intercept)          0.25      0.04    0.18    0.32 1.01
## sd(cond)                0.11      0.06    0.01    0.23 1.00
## cor(Intercept,cond)    0.01      0.35   -0.67    0.69 1.00
##                                Bulk_ESS Tail_ESS
## sd(Intercept)            1357     1933
## sd(cond)                 954      1490
## cor(Intercept,cond)      3214     2346
##
## Regression Coefficients:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept       6.00      0.06    5.88    6.13 1.00      1006     1454
## cond          -0.08      0.07   -0.22    0.06 1.00      1837     2174
##
## Further Distributional Parameters:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sigma         0.43      0.01    0.40    0.46 1.00      3169     2775
##
## Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).

```

Based on the posterior from the Gibson and Wu (2013) analysis, the posterior for the effect on the log scale is $\text{Normal}(-0.08, 0.065)$.

```

priorsgwfulllogninf <- c(set_prior("normal(6, 0.6)", class = "Intercept"),
                           set_prior("normal(-0.08,0.065)", class = "b",
                                     coef = "cond"),
                           set_prior("normal(0, 1)", class = "sd"),
                           set_prior("normal(0, 1)", class = "sigma"),
                           set_prior("lkj(2)", class = "cor"))

```

```

m_gwfulllogninf<-brm(rt~cond + (1+cond|subj)+(1+cond|item),
                       family=lognormal(), prior=priorsgwfulllogninf,
                       warmup=1000,
                       iter=2000,
                       cores=4,
                       data = df_gibsonwu2,
                       control=list(adapt_delta=0.99, max_treedepth=15))

```

```
summary(m_gwfulllogninf)
```

```

## Family: lognormal
## Links: mu = identity; sigma = identity
## Formula: rt ~ cond + (1 + cond | subj) + (1 + cond | item)
## Data: df_gibsonwu2 (Number of observations: 595)
## Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
##        total post-warmup draws = 4000
##
## Multilevel Hyperparameters:

```

```

## ~item (Number of levels: 15)
##           Estimate Est.Error l-95% CI u-95% CI Rhat
## sd(Intercept)      0.16     0.04    0.10    0.26 1.00
## sd(cond)          0.18     0.06    0.07    0.32 1.00
## cor(Intercept,cond) -0.42     0.28   -0.87    0.20 1.00
##           Bulk_ESS Tail_ESS
## sd(Intercept)      1300     2159
## sd(cond)          1293     1324
## cor(Intercept,cond) 2595     2676
##
## ~subj (Number of levels: 40)
##           Estimate Est.Error l-95% CI u-95% CI Rhat
## sd(Intercept)      0.25     0.04    0.18    0.32 1.00
## sd(cond)          0.10     0.06    0.01    0.23 1.00
## cor(Intercept,cond) 0.02     0.36   -0.71    0.71 1.00
##           Bulk_ESS Tail_ESS
## sd(Intercept)      1297     2228
## sd(cond)          999      1532
## cor(Intercept,cond) 4565     2422
##
## Regression Coefficients:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept       6.00     0.06    5.88    6.12 1.00      1386     1759
## cond          -0.08     0.04   -0.17    0.00 1.00      3869     3091
##
## Further Distributional Parameters:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sigma          0.43     0.01    0.40    0.46 1.00      4790     2888
##
## Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).

```

Next, we pool the data and fit the model again:

```

colnames(df_gibsonwu)

## [1] "subj" "item" "type" "rt"    "cond"

colnames(df_gibsonwu2)[3] <- "type"
df_gibsonwu2 <- df_gibsonwu2[,c(1,2,3,5,7)]

df_gibsonwu$subj <- paste("gw", df_gibsonwu$subj, sep="")
df_gibsonwu2$subj <- paste("gw2", df_gibsonwu2$subj, sep="")
pooled_gw <- rbind(df_gibsonwu, df_gibsonwu2)

priorsgwfulllognuninf <- c(set_prior("normal(6, 0.6)", class = "Intercept"),
                               set_prior("normal(0,1)",
                                         class = "b",
                                         coef = "cond"),
                               set_prior("normal(0, 1)", class = "sd"),
                               set_prior("normal(0, 1)", class = "sigma"),
                               set_prior("lkj(2)", class = "cor"))

m_gwfulllognpooled <- brm(rt ~ cond + (1+cond|subj)+(1+cond|item),
                            family=lognormal(), prior=priorsgwfulllognuninf,

```

```

    warmup=1000,
    iter=2000,
    cores=4,
    data = pooled_gw,
    control=list(adapt_delta=0.99, max_treedepth=15))

summary(m_gwfulllognpooled)

## Family: lognormal
## Links: mu = identity; sigma = identity
## Formula: rt ~ cond + (1 + cond | subj) + (1 + cond | item)
## Data: pooled_gw (Number of observations: 1142)
## Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
##         total post-warmup draws = 4000
##
## Multilevel Hyperparameters:
## ~item (Number of levels: 15)
##             Estimate Est.Error 1-95% CI u-95% CI Rhat
## sd(Intercept)     0.18      0.04    0.12    0.28 1.00
## sd(cond)          0.06      0.04    0.00    0.15 1.00
## cor(Intercept,cond) -0.23     0.40   -0.85    0.62 1.00
##             Bulk_ESS Tail_ESS
## sd(Intercept)     1233     2311
## sd(cond)          1515     1863
## cor(Intercept,cond) 5916     2764
##
## ~subj (Number of levels: 77)
##             Estimate Est.Error 1-95% CI u-95% CI Rhat
## sd(Intercept)     0.25      0.03    0.20    0.30 1.00
## sd(cond)          0.10      0.05    0.01    0.20 1.00
## cor(Intercept,cond) -0.32     0.30   -0.83    0.41 1.00
##             Bulk_ESS Tail_ESS
## sd(Intercept)     1469     2192
## sd(cond)          1161     1625
## cor(Intercept,cond) 4402     2194
##
## Regression Coefficients:
##             Estimate Est.Error 1-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept      6.03      0.06    5.92    6.15 1.00     1178    1617
## cond          -0.08      0.04   -0.15   -0.01 1.00     4306    2753
##
## Further Distributional Parameters:
##             Estimate Est.Error 1-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sigma          0.48      0.01    0.46    0.50 1.00     5991    3118
##
## Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).

```

We do get similar estimates by incrementally using the posterior from a previous study as a prior for a subsequent study, and from pooling the data.

Exercise 4

```
data("df_dillonE1")
dillonE1 <- df_dillonE1
```

First we set the contrast coding so that:

$$\mu_{low} = \alpha + 1 \cdot \beta$$

$$\mu_{high} = \alpha + -1 \cdot \beta$$

```
dillonE1 <- dillonE1 %>%
  mutate(int_meff = if_else(int == "low", 1, -1))

fit_dillonE1 <- brm(rt ~ int_meff + (int_meff | subj) + (int_meff | item),
  family = lognormal(),
  prior =
    c(
      prior(normal(6, 1.5), class = Intercept),
      prior(normal(0, 1), class = b),
      prior(normal(0, 1), class = sigma),
      prior(normal(0, 1), class = sd),
      prior(lkj(2), class = cor)
    ),
  data = dillonE1
)

print(fit_dillonE1)

##  Family: lognormal
##  Links: mu = identity; sigma = identity
## Formula: rt ~ int_meff + (int_meff | subj) + (int_meff | item)
##  Data: dillonE1 (Number of observations: 2855)
##  Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
##         total post-warmup draws = 4000
##
## Multilevel Hyperparameters:
## ~item (Number of levels: 48)
##             Estimate Est.Error l-95% CI u-95% CI Rhat
## sd(Intercept)     0.13     0.02     0.10     0.18 1.00
## sd(int_meff)      0.05     0.02     0.01     0.09 1.00
## cor(Intercept,int_meff)  0.25     0.28    -0.32     0.75 1.00
##                         Bulk_ESS Tail_ESS
## sd(Intercept)       1723     2807
## sd(int_meff)        897     1181
## cor(Intercept,int_meff)  2768     2517
##
## ~subj (Number of levels: 40)
##             Estimate Est.Error l-95% CI u-95% CI Rhat
## sd(Intercept)     0.28     0.04     0.22     0.36 1.00
## sd(int_meff)      0.06     0.02     0.01     0.09 1.00
## cor(Intercept,int_meff)  0.17     0.26    -0.34     0.68 1.00
##                         Bulk_ESS Tail_ESS
## sd(Intercept)       1241     2338
## sd(int_meff)        1019     1047
```

```

## cor(Intercept,int_meff)      3853      2809
##
## Regression Coefficients:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept     6.51      0.05     6.42     6.61 1.00      856     1522
## int_meff      0.03      0.02    -0.00     0.06 1.00     3937     3298
##
## Further Distributional Parameters:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sigma       0.57      0.01     0.56     0.59 1.00     6100     3014
##
## Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).
posterior_summary(fit_dillonE1, variable = "b_int_meff")

```

```

##           Estimate Est.Error   Q2.5   Q97.5
## b_int_meff  0.0287   0.0157 -0.00143  0.0601

```

There is some indication that the low interference conditions have slower rt, as predicted.

Exercise 5

```

data("df_ab")

fit_ab <- brm(probe_correct ~ lag + (lag | subj),
  family = bernoulli(),
  prior =
  c(
    prior(normal(6, 1.5), class = Intercept),
    prior(normal(0, 1), class = b),
    prior(normal(0, 1), class = sd),
    prior(lkj(2), class = cor)
  ),
  data = df_ab
)

fit_ab

```

```

## Family: bernoulli
## Links: mu = logit
## Formula: probe_correct ~ lag + (lag | subj)
## Data: df_ab (Number of observations: 2101)
## Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
##        total post-warmup draws = 4000
##
## Multilevel Hyperparameters:
## ~subj (Number of levels: 29)
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sd(Intercept)     0.60      0.14     0.37     0.90 1.00      2188
## sd(lag)          0.08      0.03     0.02     0.15 1.00      1050
## cor(Intercept,lag) 0.34      0.31    -0.27     0.88 1.00      2523
##                                     Tail_ESS
## sd(Intercept)      3015
## sd(lag)            1450

```

```

## cor(Intercept,lag)      2370
##
## Regression Coefficients:
##             Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept      0.29      0.15     0.02     0.58 1.00    2171    2394
## lag            0.02      0.02    -0.03     0.06 1.00    4345    3119
##
## Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).

```

(a) How is the accuracy of the probe identification affected by the lag? Estimate this in log-odds and percentages.

In log odds:

```
posterior_summary(fit_ab, variable = "b_lag")
```

```

##             Estimate Est.Error   Q2.5   Q97.5
## b_lag      0.016    0.0246 -0.0316  0.0646

```

From lag 0 to lag 1 the accuracy increases by

```

alpha_samples <- as_draws_df(fit_ab)$b_Intercept
beta_samples <- as_draws_df(fit_ab)$b_lag
effect_lag1_0 <- plogis(alpha_samples + 1* beta_samples) - plogis(alpha_samples)

c(mean = mean(effect_lag1_0),
quantile(effect_lag1_0, c(0.025, 0.975)))

##      mean     2.5%    97.5%
##  0.00392 -0.00781  0.01583

```

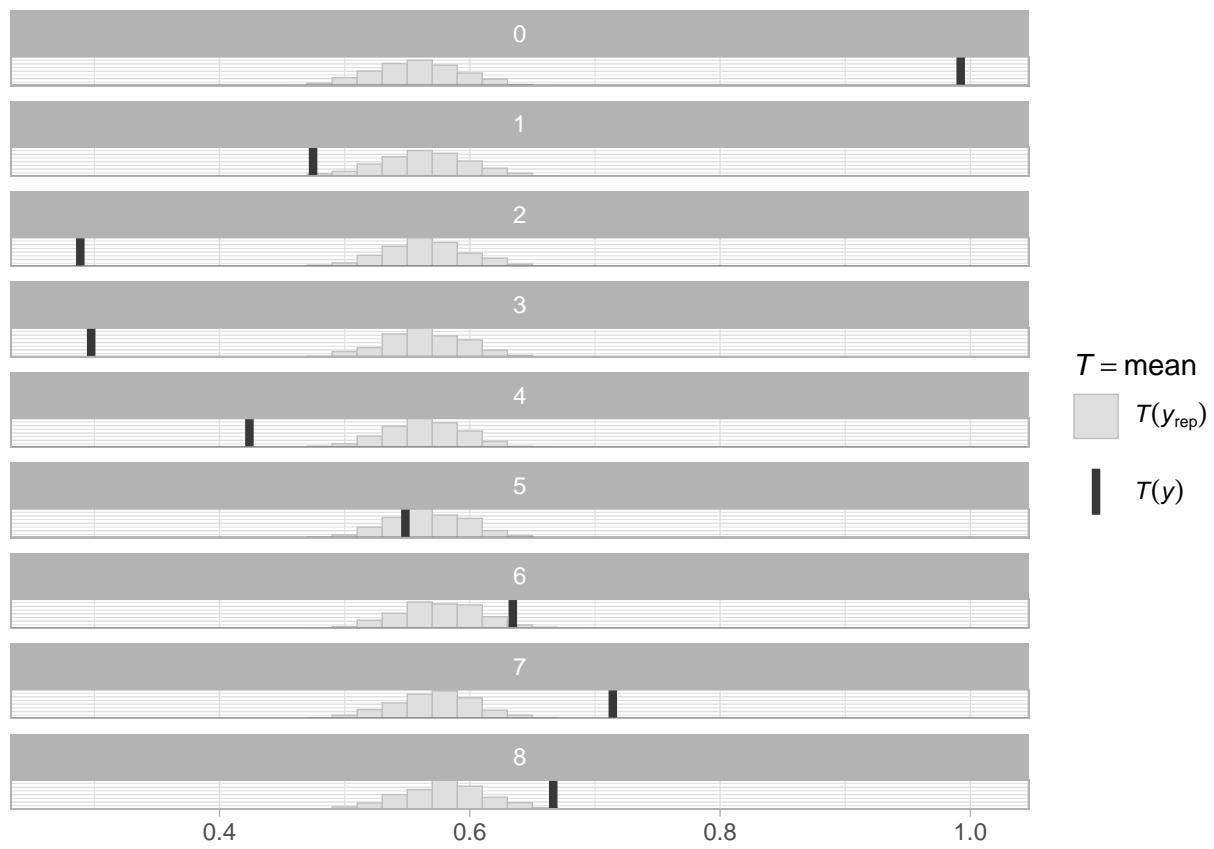
(b) Is the linear relationship justified? Use posterior predictive checks to verify this.

```

pp_check(fit_ab,
         type = "stat_grouped",
         stat = "mean",
         group = "lag",
         facet_args = list(
           ncol = 1, scales = "fixed"
         ),
         binwidth = 0.02
       )

## Using all posterior draws for ppc type 'stat_grouped' by default.

```

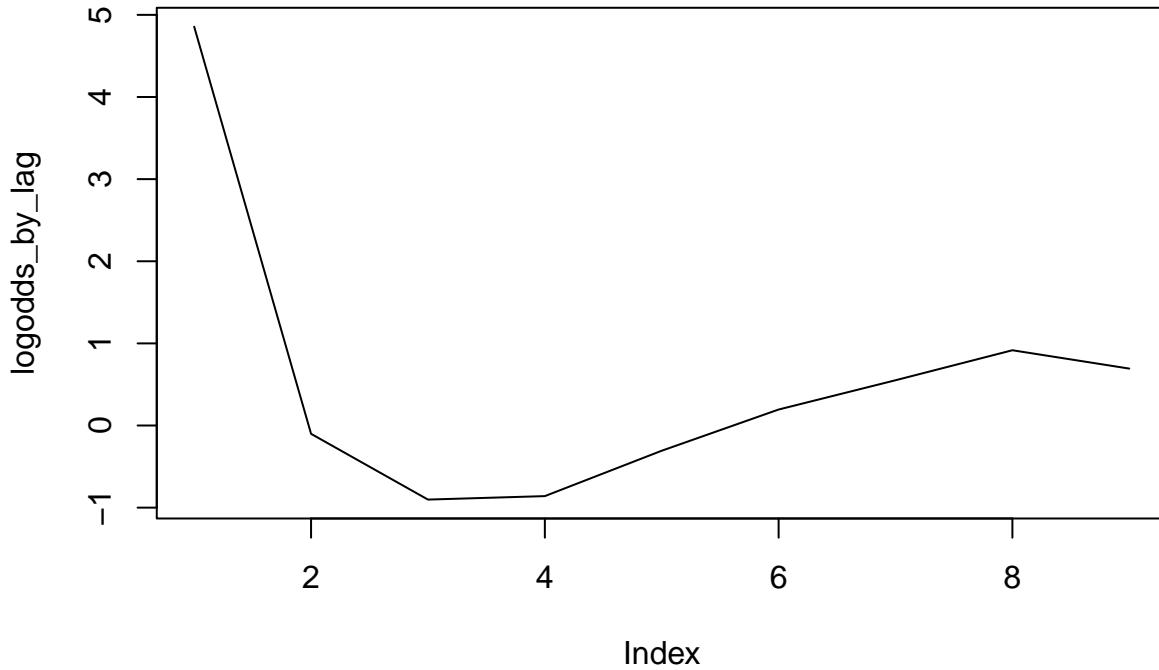


Computing the mean log odds for each lag by hand:

```
probs_by_lag<-with(df_ab,tapply(probe_correct,lag,mean))

logodds_by_lag<-qlogis(probs_by_lag)

plot(logodds_by_lag,type="l")
```



It seems that below lag 3 (lag 0, lag 1, lag 2), lag has a negative effect on the log odds of detection, and from lag $>= 3$, it reverses the lag increases the log-odds of detection. We see this in the plot (also see below): The average proportion of detections goes from 1 to 0 (in lags 0, 1, 2) in log odds space, and then it starts to increase again. A model that assumes only a monotonically increasing effect of lag can't take that into account.

This was not asked for in the exercise, but one can think about a better relationship between lag and accuracy. Fit a new model and use posterior predictive checks to verify if the fit improved.

The new model introduces a new term with an interaction to account for the U shape relationship (but the model is still linear).

- if lag ≥ 3 , the model in log odds is reduced to

$$\alpha + \beta_1 \times \text{lag}$$

- if lag < 3 , the model in log odds is

$$\alpha + \beta_1 \times \text{lag} + \beta_2 \times 1 + \beta_3 \times \text{lag} = \alpha + \beta_2 + (\beta_1 + \beta_3) \times \text{lag}$$

```
fit_ab2 <- brm(probe_correct ~ lag * I(lag < 3) + (lag * I(lag < 3) | subj),
  family = bernoulli(link= logit),
  prior =
    c(prior(normal(0, 1.5), class = Intercept),
      prior(normal(0, 0.5), class = b, coef = lag),
      prior(normal(0, 1), class = sd, coef = Intercept, group = subj),
      prior(normal(0, 1), class = sd, coef = lag, group = subj),
      prior(lkj(2), class = cor, group = subj)
    ),
  iter = 2000,
  warmup = 1000,
  control = list(adapt_delta = 0.99),
  data = df_ab
)
```

```

fit_ab2

##  Family: bernoulli
##  Links: mu = logit
## Formula: probe_correct ~ lag * I(lag < 3) + (lag * I(lag < 3) | subj)
## Data: df_ab (Number of observations: 2101)
## Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
##         total post-warmup draws = 4000
##
## Multilevel Hyperparameters:
## ~subj (Number of levels: 29)
##                               Estimate Est.Error 1-95% CI u-95% CI
## sd(Intercept)                0.98     0.20    0.63    1.42
## sd(lag)                     0.06     0.04    0.00    0.16
## sd(Ilag<3TRUE)              0.45     0.30    0.02    1.15
## sd(lag:Ilag<3TRUE)          0.80     0.23    0.40    1.31
## cor(Intercept,lag)           -0.10    0.36   -0.72    0.61
## cor(Intercept,Ilag<3TRUE)   -0.41    0.34   -0.89    0.42
## cor(lag,Ilag<3TRUE)          0.12     0.38   -0.65    0.76
## cor(Intercept,lag:Ilag<3TRUE) 0.26     0.25   -0.23    0.70
## cor(lag,lag:Ilag<3TRUE)      -0.04    0.35   -0.69    0.63
## cor(Ilag<3TRUE,lag:Ilag<3TRUE) -0.22    0.36   -0.79    0.55
##                               Rhat Bulk_ESS Tail_ESS
## sd(Intercept)                1.00    2592     2440
## sd(lag)                      1.00     485     1490
## sd(Ilag<3TRUE)               1.00    1547     1782
## sd(lag:Ilag<3TRUE)           1.00    1478     2330
## cor(Intercept,lag)            1.00    3004     3100
## cor(Intercept,Ilag<3TRUE)    1.00    3368     3353
## cor(lag,Ilag<3TRUE)           1.00    3062     2996
## cor(Intercept,lag:Ilag<3TRUE) 1.00    2341     2880
## cor(lag,lag:Ilag<3TRUE)       1.00     881     1800
## cor(Ilag<3TRUE,lag:Ilag<3TRUE) 1.00     912     1388
##
## Regression Coefficients:
##                               Estimate Est.Error 1-95% CI u-95% CI Rhat Bulk_ESS
## Intercept                  -1.92     0.28   -2.48   -1.39 1.00    2225
## lag                        0.39     0.04    0.31    0.47 1.00    5163
## Ilag<3TRUE                 4.94     0.33    4.31    5.61 1.00    4086
## lag:Ilag<3TRUE              -2.87    0.25   -3.37   -2.42 1.00    3393
##                               Tail_ESS
## Intercept                   2854
## lag                         2376
## Ilag<3TRUE                  3174
## lag:Ilag<3TRUE              3108
##
## Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).

pp_check(fit_ab2,
          type = "stat_grouped",
          stat = "mean",
          group = "lag",

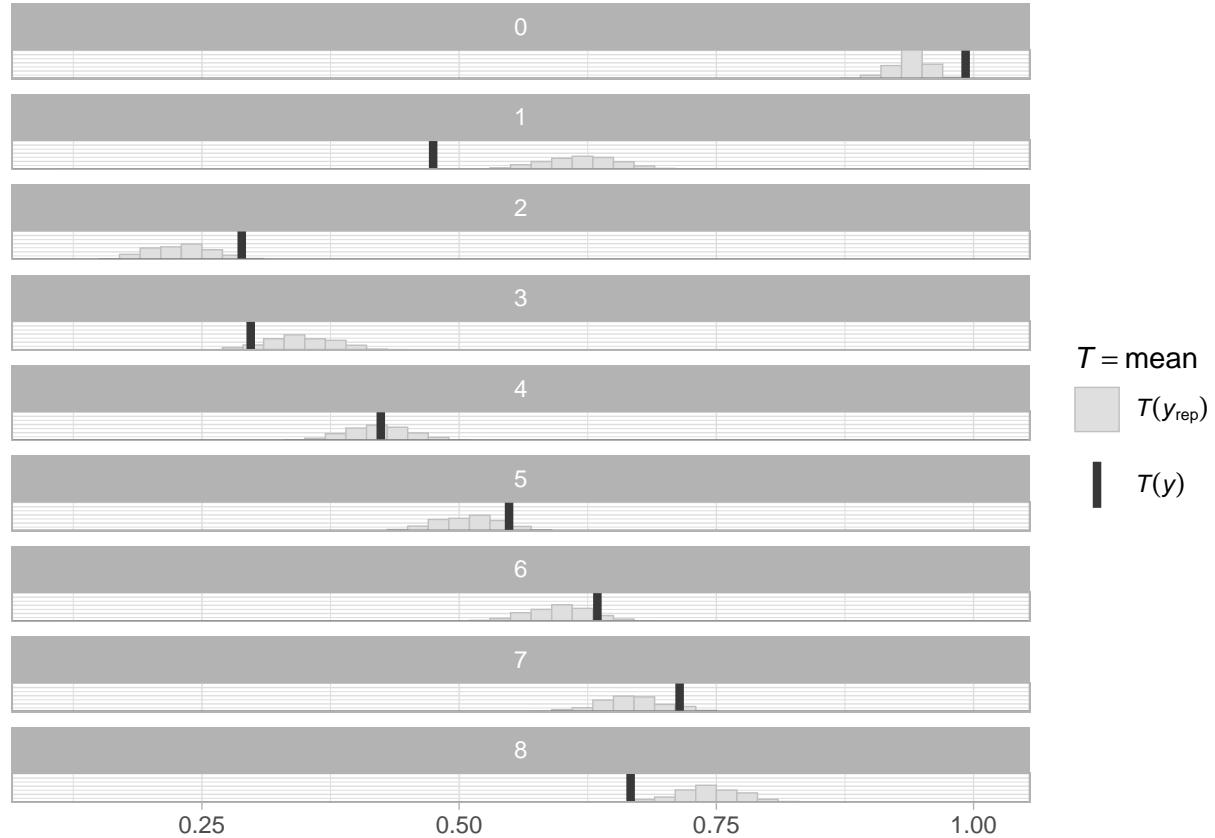
```

```

facet_args = list(
  ncol = 1, scales = "fixed"
),
binwidth = 0.02
)

```

Using all posterior draws for ppc type 'stat_grouped' by default.



This seems to account for the observed pattern a bit better.